

March 26, 2018

## **Assignment 3**

Sanchit Singhal

INF 397 – Statistical Analysis and Learning w/ Prof. Varun Rai

Spring 2018

The University of Texas at Austin

## 1) K-fold Validation

### a) Implementation

k-fold cross-validation is implemented by the dividing the full dataset of  $n$  observations into  $k$  non-overlapping groups – that usually have approximately the same size ( $n/k$ ).  $k$  number of models are then built and trained using all except one of the groups which acts as a validation set for the model one at a time. The test error is estimated by averaging the  $k$  resulting MSE estimates.

### b) Comparison to:

#### i) Validation-set approach

Although the validation-set approach is computationally less intense to execute (because you are only partitioning the data once into two sets and building/testing just 1 model), there are two major disadvantages when compared to the k-fold validation approach. The test error estimate in the validation-set approach is highly variable - it depends on which observations are chosen to be in the training set as opposed to the validation set. K-fold validation on the other hand leads to a less variable test error because it essentially uses every data point in the validation at some point as MSE is calculated for each  $k$  group and the test error is taken to be the average of all of them. Further, since the validation set approach only trains the model on a subset of the data, it will tend to perform worse than the k-fold validation approach which uses most of the data to train the model. This means that the validation-set approach can often over-estimate the test error rate on the entire dataset.

#### ii) LOOCV

The LOOCV approach can be considered a special case of the k-fold validation approach in which  $k = n$ . This approach benefits from unbiased estimates of the test error (since each model is trained on all but 1 observation). However, this approach also has two main drawbacks when compared to the k-fold validation approach. In contrast to the validation-set approach, LOOCV is computationally more expensive than the k-fold because you fit the model  $n$  times as opposed to only  $k$  times in the k-fold validation approach. This issue is additionally amplified when the model gets more complex – costing even more to build each model. Secondly, even though LOOCV can give a less biased estimate of the test error it usually has higher variance than k-fold validation. This is simply because by taking the average of  $n$  fitted models (that are trained on almost identical sets of observations), the outputs are highly correlated and thus lead to higher variance.

Usually, we should consider the bias-variance tradeoff associated when selecting our value of  $k$  – it is suggested that using  $k=5$  or  $10$  leads to test error rate estimates that suffer neither from too much bias nor variance.

```

library(class)

### Question 2

# Read CSV from working directory into R

MyData <- read.csv(file="redwine.csv", header=TRUE, sep=",")

# Create binary variable final_quality using mean

MyData$final_quality <- with(ifelse(quality>mean(quality), 1, 0), data=MyData)

# Creating dataset without original quality attribute

myvars <- names(MyData) %in% c("quality")
fulldataset <- MyData[!myvars]

## a) Logistic Regression with all the data

glm.fit <- glm(final_quality ~ fixed.acidity+volatile.acidity+citric.acid+residual.sugar+
chlorides+free.sulfur.dioxide+total.sulfur.dioxide+density+pH+sulphates+alcohol, data=fulldataset, family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = final_quality ~ fixed.acidity + volatile.acidity +
##      citric.acid + residual.sugar + chlorides + free.sulfur.dioxide +
##      total.sulfur.dioxide + density + pH + sulphates + alcohol,
##      family = binomial, data = fulldataset)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4025  -0.8387   0.3105   0.8300   2.3142
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    42.949948   79.473979   0.540  0.58890
## fixed.acidity     0.135980    0.098483   1.381  0.16736
## volatile.acidity  -3.281694    0.488214  -6.722 1.79e-11 ***
## citric.acid      -1.274347    0.562730  -2.265  0.02354 *
## residual.sugar    0.055326    0.053770   1.029  0.30351
## chlorides        -3.915713    1.569298  -2.495  0.01259 *
## free.sulfur.dioxide  0.022220    0.008236   2.698  0.00698 **
## total.sulfur.dioxide -0.016394    0.002882  -5.688 1.29e-08 ***
## density          -50.932385   81.148745  -0.628  0.53024
## pH               -0.380608    0.720203  -0.528  0.59717
## sulphates        2.795107    0.452184   6.181 6.36e-10 ***
## alcohol          0.866822    0.104190   8.320 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2209.0  on 1598  degrees of freedom

```

```

## Residual deviance: 1655.6 on 1587 degrees of freedom
## AIC: 1679.6
##
## Number of Fisher Scoring iterations: 4

## b) Logistic Regression with the validation set approach

## i) Splitting data into train and val using 80/20 split

set.seed(1)
rows <- sample(x=nrow(fulldataset), size=.80*nrow(fulldataset))
trainset <- fulldataset[rows, ]
valset <- fulldataset[-rows, ]

## ii) Logistic Regression fit with only the training set

glm2.fit <- glm(final_quality ~ fixed.acidity+volatile.acidity+citric.acid+residual.sugar
+chlorides+free.sulfur.dioxide+total.sulfur.dioxide+density+pH+sulphates+alcohol, data=trainset, family=binomial)
summary(glm2.fit)

##
## Call:
## glm(formula = final_quality ~ fixed.acidity + volatile.acidity +
##      citric.acid + residual.sugar + chlorides + free.sulfur.dioxide +
##      total.sulfur.dioxide + density + pH + sulphates + alcohol,
##      family = binomial, data = trainset)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3285  -0.8640   0.3094   0.8477   2.2338
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  19.150545  85.590617   0.224  0.82296
## fixed.acidity    0.111649   0.105894   1.054  0.29173
## volatile.acidity -2.844256   0.534364  -5.323 1.02e-07 ***
## citric.acid     -0.988874   0.630212  -1.569  0.11662
## residual.sugar   0.006617   0.061182   0.108  0.91388
## chlorides      -4.933251   1.738512  -2.838  0.00454 **
## free.sulfur.dioxide  0.016142   0.009139   1.766  0.07734 .
## total.sulfur.dioxide -0.014461   0.003119  -4.636 3.55e-06 ***
## density        -25.095517  87.419908  -0.287  0.77406
## pH             -0.921976   0.793994  -1.161  0.24557
## sulphates       2.567119   0.499801   5.136 2.80e-07 ***
## alcohol         0.872373   0.113699   7.673 1.68e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1770.7 on 1278 degrees of freedom
## Residual deviance: 1352.0 on 1267 degrees of freedom
## AIC: 1376

```

```
##
## Number of Fisher Scoring iterations: 4

## iii) Predictions for validation set: high_quality = 1, low_quality = 0

glm2.probs <- predict(glm2.fit, valset, type="response")
glm2.preds <- ifelse(glm2.probs>0.5, 1, 0)
confusion_matrix_glm2 <- table(valset$final_quality,glm2.preds)
print(confusion_matrix_glm2)

##      glm2.preds
##           0    1
## 0 105   27
## 1  45  143

## iv) Fraction of misclassified observation in validation set

Misclassified_Predictions_fraction = (confusion_matrix_glm2[1,2]+confusion_matrix_glm2[2,
1])/sum(confusion_matrix_glm2)
sprintf("Overall Fraction of Misclassified Predictions are: %f",Misclassified_Predictions
_fraction)

## [1] "Overall Fraction of Misclassified Predictions are: 0.225000"
```

### ### Question 3

#### ## a. Creating function

```
set.seed(1)
boot.fn = function(data, index) return(coef(glm(final_quality ~ fixed.acidity+volatile.ac
idity+citric.acid+residual.sugar+chlorides+free.sulfur.dioxide+total.sulfur.dioxide+densi
ty+pH+sulphates+alcohol,data = fulldataset, family = binomial, subset = index)))
```

#### ## b. Estimate standard errors

```
library(boot)
```

```
## Warning: package 'boot' was built under R version 3.4.3
```

```
boot(fulldataset, boot.fn, 5000)
```

```
##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = fulldataset, statistic = boot.fn, R = 5000)
##
##
## Bootstrap Statistics :
##      original      bias      std. error
## t1*   42.94994813 -4.1180669223  83.957974171
## t2*    0.13598034 -0.0016715528  0.105728098
## t3*   -3.28169367 -0.0458145653  0.521044067
## t4*   -1.27434734 -0.0439888532  0.599980503
## t5*    0.05532602 -0.0035509907  0.062123106
```

```
## t6*    -3.91571291 -0.1327819151  1.777481968
## t7*     0.02222037  0.0003045197  0.008274123
## t8*    -0.01639392 -0.0002640114  0.003078335
## t9*   -50.93238519  4.1679273843 85.730917155
## t10*   -0.38060751 -0.0345690167  0.737982583
## t11*    2.79510651  0.0763360253  0.522845664
## t12*    0.86682223  0.0101357328  0.107123809
```

### **## c. Comparisons of the two sets of standard errors**

*# Both sets of the estimated standard errors obtained through glm() function and  
# bootstrap function are very similar. For example, the std. error for  
# free.sulfur.dioxide was 0.008236 with glm and with the bootsrap method, it was  
# 0.008264231 (t7\*). The std. error for residual.sugar was 0.053770 with glm,  
# and 0.062763447 with bootstrap (t5\*).*